I decided to go with a CNN because I was dealing with images. In this the biggest challenge I faced was that there were so many images and the images were 200x200x3, so that lead to me having to load in the images as they were needed and not all at once. I found that in this short time I was able to get a basic set up but have not gotten as good results as I wanted.

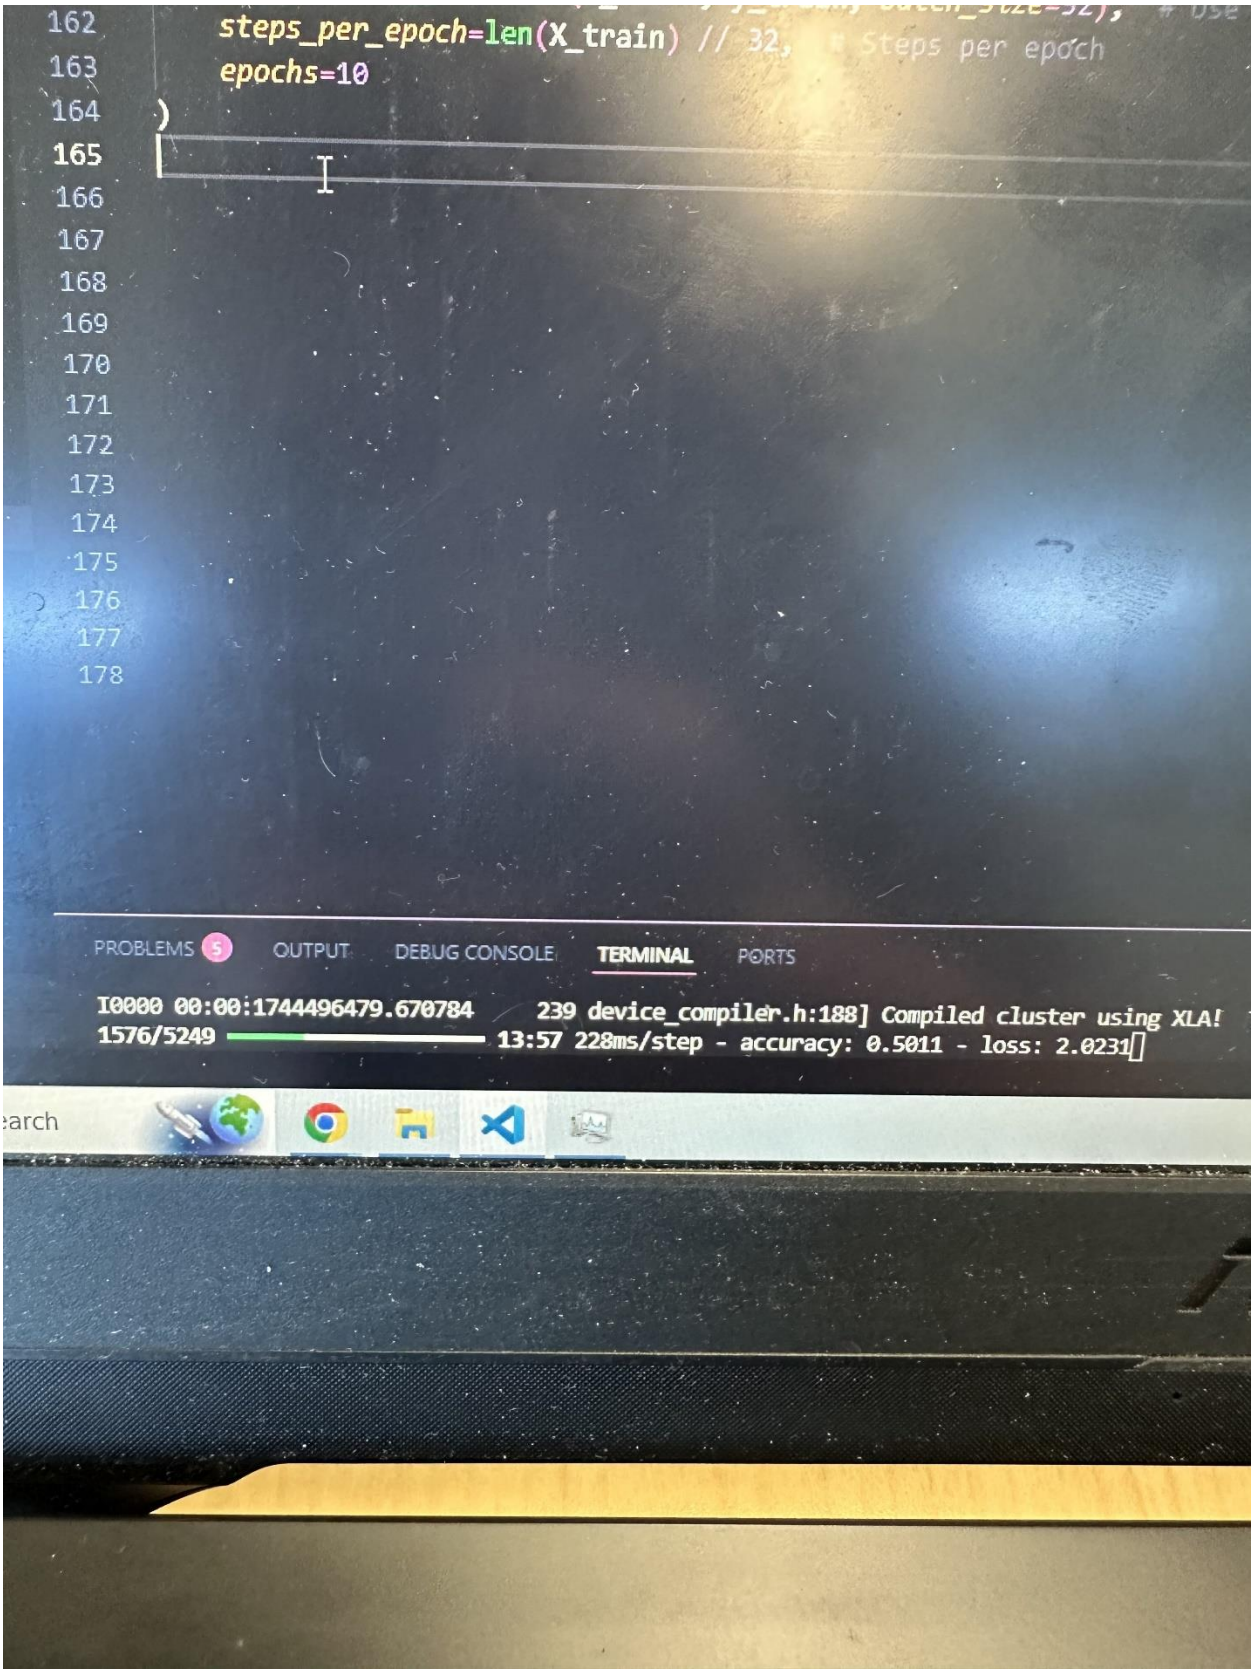
steps_per_epoch=len(X_train) // 32, # Steps per epoch
epochs=10
)
PROBLEMS 5
OUTPUT
DEBUG CONSOLE
TERMINAL
PORTS
I0000 00:00:1744496479.670784 239 device_compiler.h:188] Compiled cluster using XLA!
1576/5249 13:57 228ms/step - accuracy: 0.5011 - loss: 2.0231

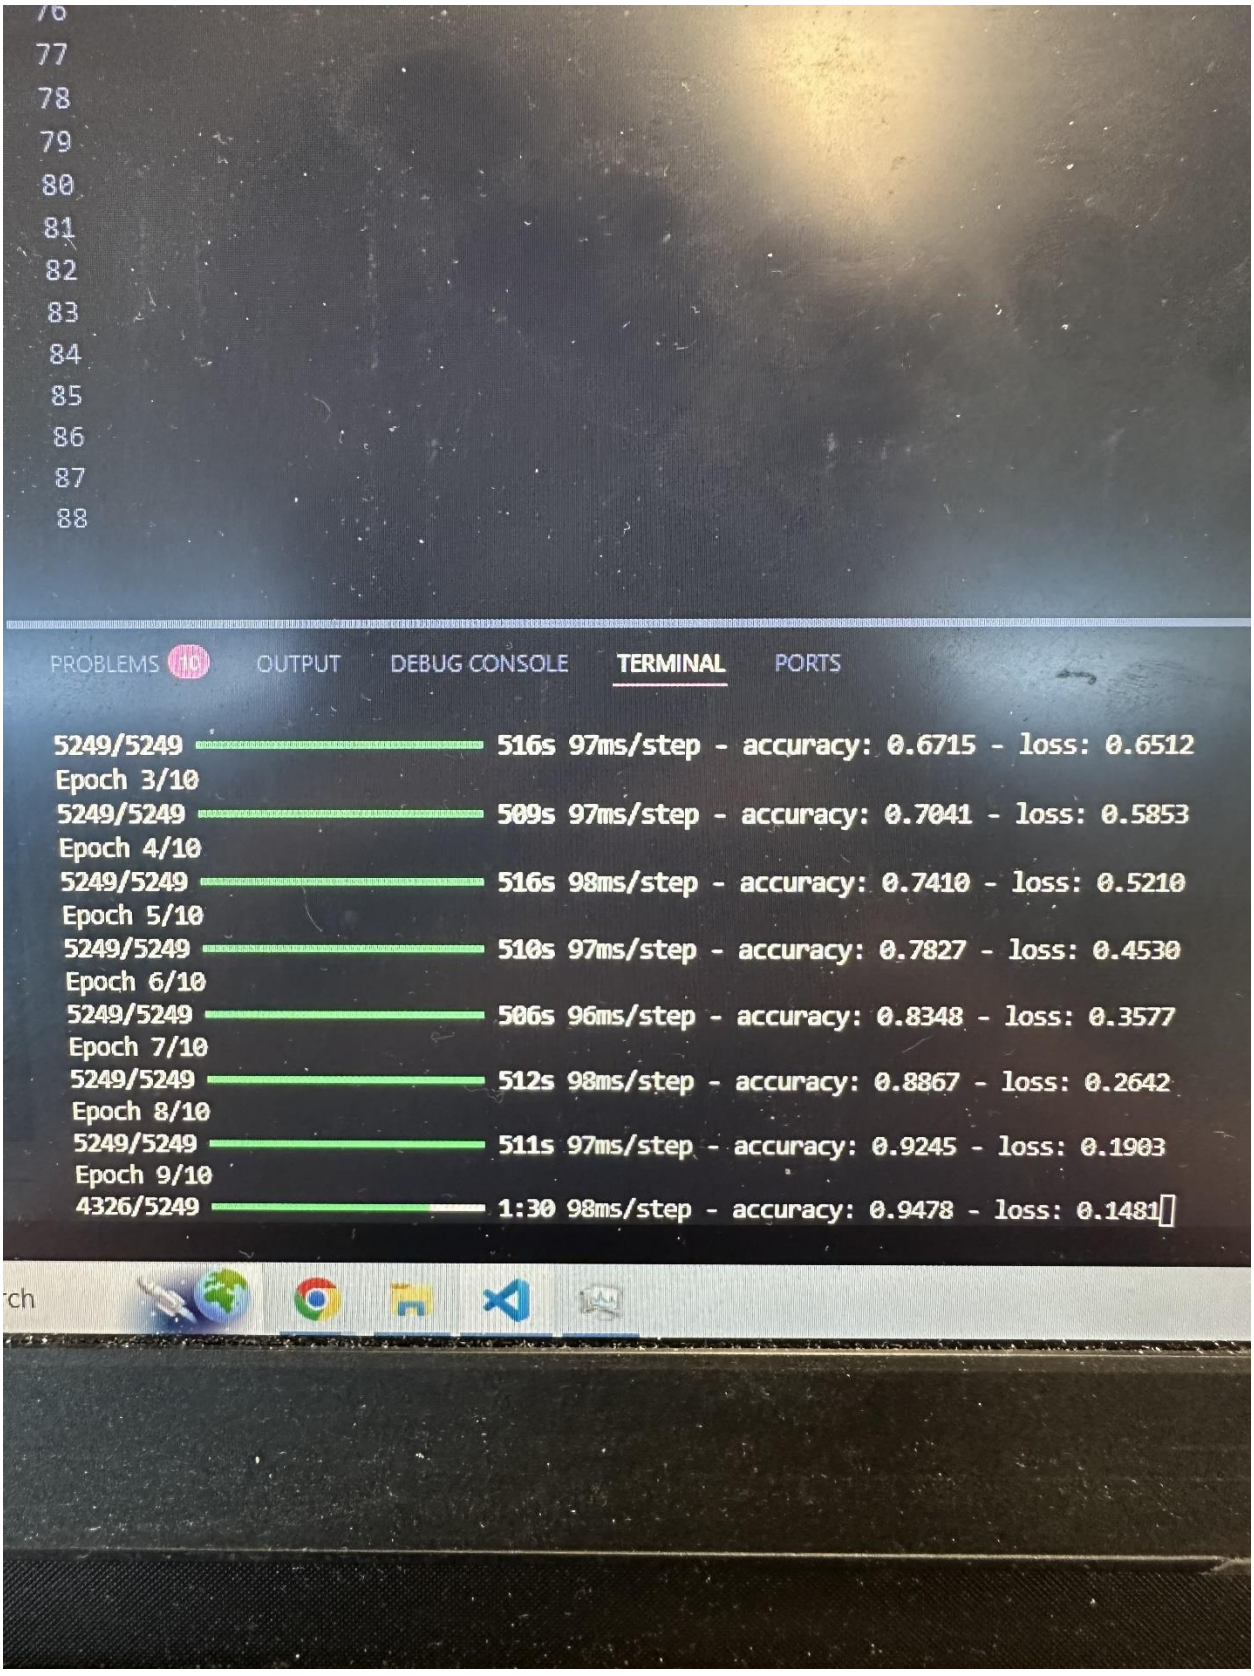
77
78
79
80
81
82
83
84
85
86
87
88
PROBLEMS 10
OUTPUT
DEBUG CONSOLE
TERMINAL
PORTS
5249/5249 516s 97ms/step - accuracy: 0.6715 - loss: 0.6512
Epoch 3/10
5249/5249 509s 97ms/step - accuracy: 0.7041 - loss: 0.5853
Epoch 4/10
5249/5249 516s 98ms/step - accuracy: 0.7410 - loss: 0.5210
Epoch 5/10
5249/5249 510s 97ms/step - accuracy: 0.7827 - loss: 0.4530
Epoch 6/10
5249/5249 506s 96ms/step - accuracy: 0.8348 - loss: 0.3577
Epoch 7/10
5249/5249 512s 98ms/step - accuracy: 0.8867 - loss: 0.2642
Epoch 8/10
5249/5249 511s 97ms/step - accuracy: 0.9245 - loss: 0.1903
Epoch 9/10
4326/5249 1:30 98ms/step - accuracy: 0.9478 - loss: 0.1481